

Backdooring and Poisoning Neural Networks with Image-Scaling Attacks

Erwin Quiring and Konrad Rieck
Technische Universität Braunschweig, Germany

Abstract—Backdoors and poisoning attacks are a major threat to the security of machine-learning and vision systems. Often, however, these attacks leave visible artifacts in the images that can be visually detected and weaken the efficacy of the attacks. In this paper, we propose a novel strategy for hiding backdoor and poisoning attacks. Our approach builds on a recent class of attacks against image scaling. These attacks enable manipulating images such that they change their content when scaled to a specific resolution. By combining poisoning and image-scaling attacks, we can conceal the trigger of backdoors as well as hide the overlays of clean-label poisoning. Furthermore, we consider the detection of image-scaling attacks and derive an adaptive attack. In an empirical evaluation, we demonstrate the effectiveness of our strategy. First, we show that backdoors and poisoning work equally well when combined with image-scaling attacks. Second, we demonstrate that current detection defenses against image-scaling attacks are insufficient to uncover our manipulations. Overall, our work provides a novel means for hiding traces of manipulations, being applicable to different poisoning approaches.

I. INTRODUCTION

Machine Learning is nowadays used in various security-critical applications that range from intrusion detection and medical systems to autonomous cars. Despite remarkable results, research on the security of machine learning has revealed various possible attacks. A considerable threat are poisoning attacks during the training process [e.g. 1, 5, 8]. Deep learning applications usually require a large number of training instances, so that there is a risk of an insider carefully manipulating a portion of the training data. Moreover, the training process can be outsourced either due to a lack of expertise in deep learning or due to missing computational resources to train large networks—again giving the chance to manipulate training data and the model.

In the context of deep learning, recent research has demonstrated that neural networks can be modified to return targeted responses without an impact on their behavior for benign inputs. An adversary, for instance, can insert a pattern in some training images of a particular target class, so that the network learns to associate the pattern with this class. If the pattern is added to arbitrary images, the network returns the target class. However, a major drawback of most attacks is the visibility of data manipulations either at training or test time [5, 8]. The attack is thus revealed if human beings audit the respective images.

Published at IEEE Deep Learning and Security Workshop (DLS) 2020, co-located with the 41st IEEE Symposium on Security and Privacy (S&P)

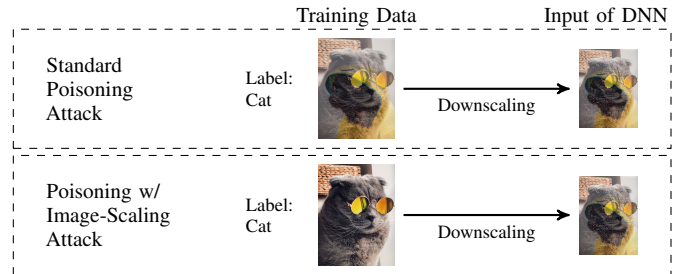


Fig. 1. Example of a clean-label poisoning attack [12]: a neural network learns to classify a dog as cat by blending the dog with multiple cat images. Image-scaling attacks allow more insidious poisoning attacks. The dog as manipulation is not visible in the training data and appears only after downscaling.

Xiao et al. [20] have recently presented a novel attack vulnerability in the data preprocessing of typical machine learning pipelines. An adversary can slightly manipulate an image, such that an *image scaling* algorithm produces a novel and unrelated image in the network’s input dimensions. The attack exploits that images are typically larger than the input dimensions and thus need to be scaled.

This novel attack directly addresses the shortcomings of most poisoning attacks by allowing an adversary to conceal data manipulations. As an example, Figure 1 shows a clean-label poisoning attack [12] on the popular TensorFlow library. The network will learn to classify the dog as cat if this dog is repeatedly inserted into varying images showing cats during training. In the attack’s standard version, the slight manipulation of the training image is still noticeable. Yet, image-scaling attacks conceal the manipulation of the training data effectively. The dog appears only in the downsampled image which is finally used by the neural network.

This paper provides the first analysis on the combination of data poisoning and image-scaling attacks. Our findings show that an adversary can significantly conceal image manipulations of current backdoor attacks [5] and clean-label attacks [12] without an impact on their overall attack success rate. Moreover, we demonstrate that defenses—designed to detect image-scaling attacks—fail in the poisoning scenario. We examine the histogram- and color-scattering-based detection as proposed by Xiao et al. [20]. In an empirical evaluation, we show that both defenses cannot detect backdoor attacks due to bounded, local changes. We further derive a novel adaptive attack that significantly reduces the performance of both defenses in the clean-label setting. All in all, our findings indicate a need for novel, robust detection defenses against image-scaling attacks.

Contributions. In summary, we make the following contributions in this paper:

- *Combination of data poisoning and image-scaling attacks.* We provide the first analysis on poisoning attacks that are combined with image-scaling attacks. We discuss two realistic threat models and consider backdoor attacks as well as clean-label poisoning attacks.
- *Evaluation of defenses* We evaluate current detection methods against image-scaling attacks and show that backdoor attacks cannot be detected.
- *Adaptive Attack.* We derive a novel variant of image-scaling attack that reduces the detection rate of current scaling defenses. Our evaluation shows that clean-label attacks cannot be reliably detected anymore.

The remainder of this paper is organized as follows: Section II reviews the background of data poisoning and image-scaling attacks. Section III examines their combination with the respective threat scenarios and our adaptive attack. Section IV provides an empirical evaluation of attacks and defenses. Section V and VI present limitations and related work, respectively. Section VII concludes the paper.

II. BACKGROUND

Let us start by briefly examining poisoning and image-scaling attacks on machine learning. Both attacks operate at different stages in a typical machine learning pipeline and allow more powerful attacks when combined, as we will show in the remainder of this work.

A. Poisoning Attacks in Machine Learning

In machine learning, the training process is one of the most critical steps due to the impact on all subsequent applications. At this stage, poisoning attacks allow an adversary to change the overall model behavior [e.g. 1, 6] or to obtain targeted responses for specific inputs [e.g. 5, 8, 12] by manipulating the training data or learning model. Such attacks need to be considered whenever the training process is outsourced or an adversary has direct access to the data or model as insider [15]. Moreover, a possible manipulation needs to be considered if a learning model is continuously updated with external data.

In this work, we focus on poisoning attacks against deep neural networks where the adversary manipulates the training data to obtain targeted predictions at test time. While particularly effective with a few changed training instances, most methods have the major shortcoming that the manipulation is visible [e.g. 5, 8]. As a result, the attack can be easily uncovered if the dataset is, for instance, audited by human beings. We present two representative poisoning attacks in Section III and show that they can be easily combined with image-scaling attacks to conceal manipulations significantly.

B. Image-Scaling Attacks

The preprocessing stage in a typical machine learning pipeline is another critical point, surprisingly overlooked by previous work so far. Xiao et al. [20] have recently identified

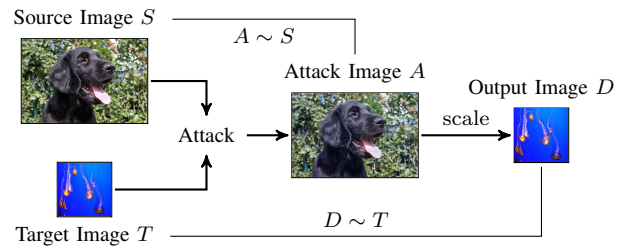


Fig. 2. Principle of image-scaling attacks: An adversary computes A such that it looks like S but downscales to T .

an attack possibility in the scaling routine of common machine learning frameworks. The attack exploits that most learning-based models expect a fixed-size input, such as 224×224 pixels for VGG19 and 299×299 pixels for InceptionV3 [14, 16]. As images are usually larger than the input dimensions of learning models, a downscaling operation as preprocessing stage is mandatory. In this case, an adversary can slightly modify an image such that it changes its content after downscaling. She can thus create targeted inputs for a neural network being invisible in the original resolution before, as exemplified by Figure 2.

Attack. In particular, the adversary slightly modifies a source image S such that the resulting attack image $A = S + \Delta$ matches a target image T after scaling. The attack can be modeled as the following quadratic optimization problem:

$$\min(\|\Delta\|_2^2) \text{ s.t. } \|\text{scale}(S + \Delta) - T\|_\infty \leq \epsilon. \quad (1)$$

Moreover, each pixel of A needs to stay in the range of $[0, 255]$ for 8-bit images. Note that an image-scaling attack is successful only if the following two goals are fulfilled:

- (O1) The downscaled output D of the attack image A is close to the target image: $D \sim T$.
- (O2) The attack image A needs to be indistinguishable from the source image: $A \sim S$.

For a detailed root-cause analysis of image-scaling attacks, we refer the reader to Quiring et al. [10].

Detection. Two methods have been proposed to *detect* image-scaling attacks [20], that is, decide if an image was manipulated to cause another result after downscaling. Both rest on the following idea, exemplified by Figure 3: The downscaling of A creates a novel image D which is unrelated to the original content from A . If we upscale D back to its original resolution, we can compare A and the upscaled version A' . In the case of an attack, both images will be different to each other.

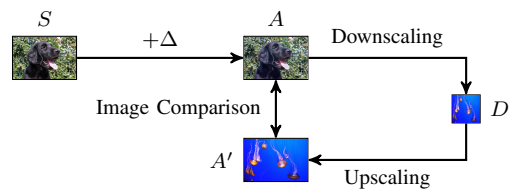


Fig. 3. Defense based on down- and upscaling with image comparison. The downscaled version D of A is upscaled again and compared with A .

The first method uses an *intensity histogram* that counts the number of pixels for each value in the dynamic range of an image. To this end, a color image is converted into a grayscale image before computing its histogram. The result is a 256 dimensional vector v^h for 8-bit images. The attack detection is now based on the cosine similarity between A and A' : $s^h = \cos(v_1^h, v_2^h)$. A low score indicates an attack, as the distribution of both inputs do not match to each other.

The second method based on *color scattering* considers spatial relations in an image. The color image is again converted to grayscale, and the average distance to the image center over all pixels with the same value is calculated, respectively. This forms a 256 dimensional vector v^s . The respective vectors from A and A' are also compared by using the cosine similarity.

We finally note that defenses exist to *prevent* image-scaling attacks [see 10]. In contrast to detection, prevention blocks the attack from the beginning, but would not uncover that the dataset was manipulated, which is the focus in this work.

III. DATA POISONING USING IMAGE-SCALING

Equipped with an understanding of data poisoning and image-scaling attacks, we are ready to examine novel attacks. The adversary’s goal is that the model returns targeted predictions for specially crafted inputs while behaving normally for benign inputs. As image-scaling attacks provide a new means for creating novel content after downscaling, they are a perfect candidate to create less visible poisoning attacks. We start by describing two plausible threat models and continue with a description of two attack variants.

A. Two Realistic Threat Models

Stealthiness during test time. It is common practice to outsource the training of large deep neural networks, either due to the lack of computational resources or due to missing expertise. In this scenario, an adversary can arbitrarily change the training data (or the model), as long as the returned model has the expected accuracy and architecture. The application of image-scaling attacks to hide changes is here not necessary. However, common backdoor attacks add visible triggers in test time instances to obtain a targeted prediction [e.g. 5, 8]. If such instances are examined, a visible backdoor, for instance, would directly reveal that a model has been tampered. These attacks can thus benefit from image-scaling attacks at test time.

Stealthiness during training time. In the second scenario, the adversary has only access to the training data, but the training process or model cannot be manipulated. This scenario is particularly relevant with insiders who already have privileged access within a company [15]. Most poisoning attacks leave visible traces in the training data, so that the attack is detectable if audited by human beings. Consequently, image-scaling attacks are also useful for these scenarios.

Finally, the application of image-scaling attacks requires (a) knowledge of the used scaling algorithm and (b) the input size of the neural network. Their knowledge is plausible to assume if the attacker is an insider or trains the model herself.

B. Enhanced Poisoning Attacks

We study two representative poisoning attacks against deep neural networks: *backdoor* and *clean-label* attacks. Both enable us to examine different approaches to manipulate images and their impact on image-scaling attacks and defenses.

Backdoor attack. As first attack, we use the *BadNets* backdoor method from Gu et al. [5]. The adversary chooses a target label and a small, bounded backdoor pattern. This pattern is added to a limited number of training images and the respective label is changed to the target label. In this way, the classifier starts to associate this pattern with the target class.

We consider both threat models for the attack. As first variant, the adversary hides the poisoning on test time instances only. Thus, we use the *BadNets* method in its classic variant during the training process. At test time, the adversary applies an image-scaling attack. The original image without backdoor represents the source image S , its version with the backdoor in the network’s input dimensions is the target image T . By solving Eq. (1), the adversary obtains the attack image A that is passed to the learning system. The pattern is only present after downscaling, so that an adversary can effectively disguise the neural network’s backdoor.

In addition, we study the threat scenario where the adversary hides the modifications at training time. We use the same attack principle as before, but apply the image-scaling attack for the backdoored training images as well. This scenario is especially relevant if the backdoor is implemented in the physical world, e.g. on road signs. The trigger can be disguised in the training data by using image-scaling attacks, and easily activated in the physical world at test time (without a scaling attack).

Clean-label poisoning attack. As second attack, we consider the poisoning attack at training time as proposed by Shafahi et al. [12]. The attack does not change the label of the modified training instances. As a result, this poisoning strategy becomes more powerful in combination with image-scaling attacks: The manipulated images keep their correct class label and show no obvious traces of manipulation.

In particular, the adversary’s objective is that the model classifies a specific and unmodified test set instance Z as a chosen target class c_t . To this end, the adversary chooses a set of images X_i from c_t . Similar to watermarking, she embeds a low-opacity version of Z into each image:

$$X'_i = \alpha \cdot Z + (1 - \alpha) \cdot X_i. \quad (2)$$

If the parameter α , for instance, is set to 0.3, features of Z are blended into X_i while the manipulation is less visible. For an image-scaling attack, the adversary chooses X_i as source image S , and creates X'_i as respective target image T in the network’s input dimensions. The computed attack image A serves as training image then. The changed images are finally added to the training set together with their correct label c_t . As a result, the classifier learns to associate Z with c_t . At test time, Z can be passed to the learning system without any changes and is classified as c_t . This attack enables us to study

the detection of image-scaling attacks, if the entire image is slightly changed instead of adding a small and bounded trigger.

C. Adaptive Image-Scaling Attack

To hide poisoned images from detection, we additionally introduce a new variant of image-scaling attack. In particular, it targets the histogram-based defense, but is also effective against the color-scattering-based approach.

The difficult part is to create an attack image A that changes its appearance to T after downscaling, but has a similar histogram if upscaled again, denoted as A' . To this end, we use the following strategy: we upscale the target image T and perform a histogram matching to the source image S . After slightly denoising the result to make the adjusted histogram smoother, we downscale the adapted image which gives us T' . We finally mount the image-scaling attack with S as source image and T' as target. Although the content changes after down- and upscaling, the histogram remains similar.

Figures 4(a) and 4(b) show an example with the histograms of A and A' for the original attack and our adapted attack, respectively. Our adaptive attack enables aligning the histograms of A and A' although both are visually different to each other, as depicted by Figures 4(f) and (h). Moreover, the visual differences to the original attack are marginal.

However, the previous example also underlines that we do not obtain an exact histogram matching. The attack chances increase if source- and target image already have a similar color tone. Therefore, we let the adversary select the most suitable images for the attack. She mounts the attack on a larger number of adapted images and select those with the highest score.

IV. EVALUATION

We continue with an empirical evaluation and perform the following experiments:

- 1) *Poisoning & image-scaling attacks.* We first demonstrate that poisoning attacks benefit from image-scaling attacks. The attack performance remains constant while the data manipulation is hard to notice.
- 2) *Detection defenses.* We demonstrate that currently proposed defenses against image-scaling attacks cannot detect backdoor attacks.
- 3) *Adaptive attack.* We also show that clean-label attacks cannot be detected if our new adaptive attack is applied to the manipulated images.

A. Dataset & Setup

For our evaluation, we use the CIFAR-10 dataset [7]. Its respective default training set is further separated into a training (40,000 images) and validation set (10,000 images) that are used for model training. We choose the model architecture from Carlini and Wagner [3] which is commonly used in the adversarial learning literature. The model expects input images of size $32 \times 32 \times 3$. This simple configuration of dataset and model allows us to train a neural network for a variety of different attack configurations in feasible time.

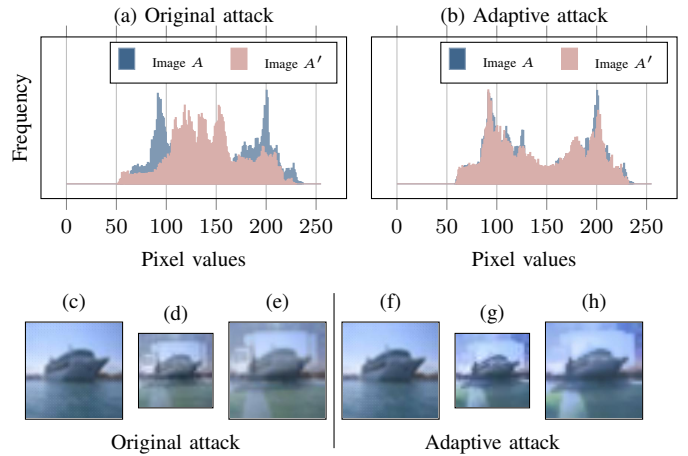


Fig. 4. Example of our adaptive image-scaling attack. Plot (a) and (b) show the compared histograms for the original and our adaptive attack. Plot (c) and (f) show A by using the original attack and our adaptive version, respectively. Plot (d) and (g) show their respective downscaled version as input for the neural network, (e) and (h) the respective upscaled version A' .

We implement the image-scaling attack in the strong variant as proposed by Xiao et al. [20], and set $\epsilon = 1.0$ in Eq. (1). We use TensorFlow (version 1.13) and report results for bilinear scaling, which is the default algorithm in TensorFlow. Due to our controlled scaling ratio, other scaling algorithms work identically and thus are omitted [see 10].

To evaluate image-scaling attacks realistically, we need source images in a higher resolution. To this end, we consider common scaling ratios from the ImageNet dataset [11]. Its images are considerably larger than the input sizes of popular models for this dataset. VGG19, for instance, expects images with size $224 \times 224 \times 3$ [14]. Based on these results, we upscale the CIFAR-10 images to a size of $256 \times 256 \times 3$ by using OpenCV’s Lanczos algorithm. This avoids side effects if the same algorithm is used for upscaling and downscaling during an image-scaling attack and model training¹.

B. Backdoor Attacks

Our first experiment tests whether image-scaling attacks can effectively conceal backdoors. For a given target class, we embed a filled black square in the lower-left corner as backdoor into training images. We perform the experiments for each class, respectively. To assess the impact of backdooring on benign inputs, we evaluate the accuracy on the unmodified CIFAR-10 test set. When evaluating backdoors on the test set, we exclude images from the target class. We report averaged results over the ten target classes. For each experiment, a baseline is added if the backdoor attack is applied without using an image-scaling attack.

Attack performance. Figure 5 presents the success rate of the original attack for a varying number of backdoored training images. The adversary can successfully control the prediction

¹If we use the same algorithm, we obtained even better results against image-scaling defenses, which might not be realistic with real-world images.

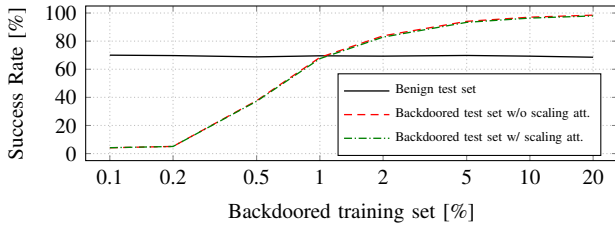


Fig. 5. Backdoor attacks: Percentage of obtained target classes on the backdoored test set, with and without image-scaling attacks for hiding the backdoor. Scaling attacks have no negative impact on the attack’s success rate.




Fig. 6. Backdoor attack examples. The first and second row result in the third row after downscaling. However, the second row relies on image-scaling attacks and better hides the backdoor trigger.

by embedding a backdoor on test time instances. If 5% of the training data are changed, she obtains an almost perfect attack result. The test set accuracy with unmodified images does not change considerably.

The application of image-scaling attacks on the backdoored test time instances has no negative impact on the success rate. At the same, the adversary can considerably hide the backdoor in contrast to the original attack, as Figure 6 shows. Although the backdoor’s high contrast with neighboring pixels and its locality creates rather unusual noise in the backdoor area, the detection is hard if only quickly audited.

In addition, we evaluate the variant where the image-scaling attack is also applied on the *training* data to hide the backdoor pattern. We obtain identical results to the previous scenario regarding the success rate and visibility of the backdoor pattern. In summary, image-scaling attacks considerably raise the bar to detect backdoors.

Detection of image-scaling attacks. Another relevant question concerns the reliable detection of image-scaling attacks (see Section II-B). Figure 7 depicts ROC curves for the histogram-based and color-scattering-based defense when backdoors are inserted at test time or training time.

For both threat scenarios, the defenses fail to detect image-scaling attacks. A closer analysis reveals that the backdoor manipulation is too small compared to the overall image size. Thus, down- and upscaling creates an image version that still corresponds to the respective input. We conclude that a reliable attack detection is thus not possible if small and bounded parts of an image are changed only.

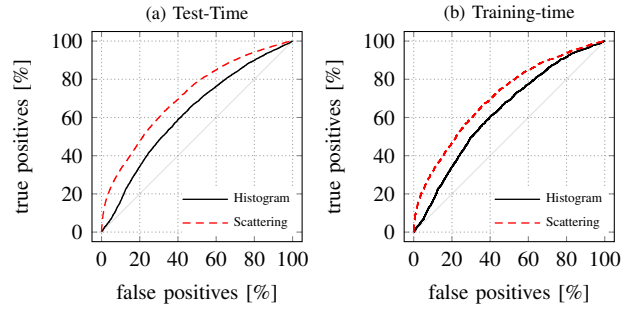


Fig. 7. Defenses against backdoor attacks: ROC curves of histogram-based and color scattering-based method. Both do not reliably detect image-scaling attacks that hide backdoors.

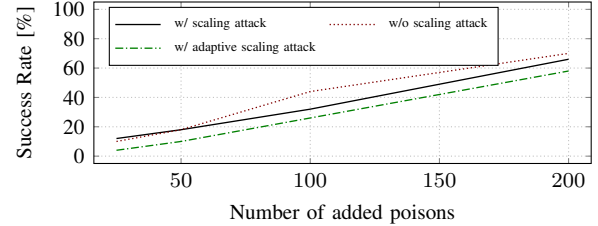


Fig. 8. Clean-label attacks: Efficiency of attack in controlling the prediction with and without image-scaling attacks, and our adaptive variant.

C. Clean-Label Poisoning Attack

We proceed with the clean-label attack from Section III-B, following the experimental setup from Shafahi et al. [12]. We test 50 randomly selected target-source class pairs (c_t, c_z) where c_z denotes the original class of Z , c_t the adversary’s target class. For each pair, we choose a random target instance Z and vary the number of modified images X_i . Again, a baseline is added where no image-scaling attack is applied on X_i . For the embedding, we set $\alpha = 0.3$.

Attack performance. Figure 8 presents the success rate of the attack with respect to the number of modified images. The adversary can significantly control the prediction for Z . The success rate increases with a larger number of modified images that are added to the training set, and corresponds to results as reported by Shafahi et al. [12].

Image-scaling attacks have only a slight impact on the success rate of the poisoning attack. At the same time, the attacker can conceal the added content of Z effectively, as exemplified by Figure 9. The 4th row emphasizes that the added content is not visible in the novel images used for training, while Z is visible for the original attack.

As opposed to the backdoor attack from the previous section, the added content Z from the clean-label attack is not noticeable even under a closer analysis. As the whole image is partly changed, the manipulation becomes an imperceptible noise pattern. We conclude that poisoning attacks can benefit from image-scaling attacks the most if the manipulation is a weaker signal, distributed over a larger area in an image.

Detection of image-scaling attacks. Figure 10(a) depicts ROC curves for the defenses. Only the histogram-based method



Fig. 9. Clean-label attack examples. The 3rd and 4th row result in the 5th row after downscaling. Image-scaling attacks can effectively hide the manipulation (4th row), which would be visible in the original poisoning attack (3rd row).

can reliably detect attacks. At 1% false positives, 94.5% of manipulated images are correctly marked as attack. The color-scattering-based approach detects only 48.2% at 1% false positives. In contrast to backdoor attacks, both defenses can more reliably spot the manipulations by the clean-label attack. As the whole image is slightly changed, the difference between the attack image and its down- and up-scaled version increases—enabling the detection.

D. Adaptive attack

We finally demonstrate that an adversary can use an adaptive strategy against both defenses to lower their detection rate. Figure 10(b) presents ROC curves for our adaptive image-scaling attack in the clean-label scenario. Our attack significantly lowers the detection rate. At the same time, the overall success rate of the attack is only slightly affected (see Figure 8). We contribute this to the histogram matching, so that parts of Z are slightly weaker embedded, especially for very dark or highly saturated images. Overall, we conclude that an adversary can circumvent current detection methods by adjusting histograms.

V. LIMITATIONS

Our findings demonstrate the benefit of image-scaling attacks for poisoning and the need to find novel detection defenses. Nonetheless, our analysis has limitations that we discuss in the following. First, we consider defenses against image-scaling attacks only. Direct defenses against data poisoning [e.g. 19] are another possible line of defense that would need to be used after downscaling. The design of such defenses is an ongoing research problem [17, 19] and beyond the scope of this work. Furthermore, we apply a simple backdoor technique by adding a filled box into training images. We do not optimize regarding shapes or the model architecture, resulting in a relatively high amount of manipulated training data. Our goal is rather to draw first conclusions about the utility of image-scaling attacks for backdooring. As scaling attacks are agnostic to the model and poisoning attack, other backdoor techniques are also applicable whenever a manipulation needs to be concealed.

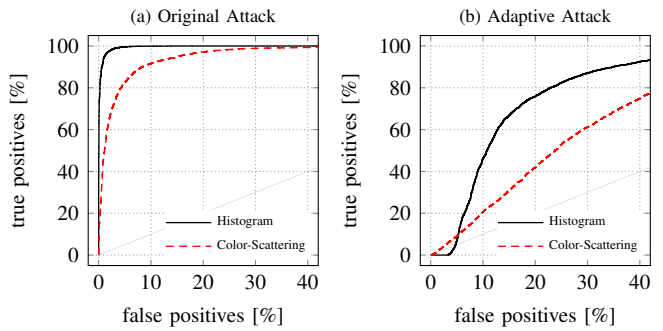


Fig. 10. Defenses against clean-label attacks: ROC curves of histogram-based and color scattering-based method with the original and adaptive attack.

VI. RELATED WORK

The secure application of machine learning requires considering various attacks along a typical workflow. Regarding the order of the targeted step, attacks can be categorized into the following classes: membership inference [e.g., 13], poisoning attacks [e.g., 1, 5, 8], evasion- and perturbation attacks [e.g., 2, 3, 9], as well as model stealing [e.g., 18].

In this work, we focus on poisoning attacks that manipulate the training data so that the learning model returns targeted responses with adversarial inputs only while behaving normally for benign inputs. Two attack variants are backdoor and clean-label poisoning attacks, differing in the amount of necessary data changes, visibility or robustness with transfer learning [e.g. 4, 5, 8, 12, 21]. We consider the following two rather simple, but representative approaches: The BadNets method [5] inserts a small, bounded pattern into images as backdoor, while the clean-label poisoning attack from Shafahi et al. [12] slightly changes the whole image to add a poison. Both provide first insights about the applicability of image-scaling attacks for data poisoning.

Concurrently, Quiring et al. [10] comprehensively analyze image-scaling attacks by identifying the root-cause and examining defenses for *prevention*. Our work here extends this line of research on image-scaling attacks by analyzing the poisoning application and *detection* defenses. While prevention stops any attack, detection uncovers that an attack is going on. Our findings here underline the need for novel detection approaches.

VII. CONCLUSION

This work demonstrates that image-scaling attacks can be leveraged to hide data manipulations for poisoning attacks. We consider two representative approaches: a backdoor attack [5] and a clean-label poisoning attack [12]. Our evaluation shows that the adversary can conceal manipulations more effectively without impact on the overall success rate of her poisoning attack. We find that image-scaling attacks can create almost invisible poisoned instances if a slight manipulation is spread over a larger area of the input.

Furthermore, our work raises the need for novel detection defenses against image-scaling attacks. Local and bounded changes—as done for backdoors—are not detected at all. The detection if the whole image is changed can be circumvented by using our proposed adaptive image-scaling attack variant.

AVAILABILITY

We make our dataset and code publicly available at <http://scaling-attacks.net> to encourage further research on poisoning attacks and image-scaling attacks.

ACKNOWLEDGMENT

The authors gratefully acknowledge funding by the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) under Germany's Excellence Strategy - EXC 2092 CASA - 390781972 and the research grant RI 2469/3-1, as well as by the German Ministry for Education and Research as BIFOLD - Berlin Institute for the Foundations of Learning and Data (ref. 01IS18025A and ref 01IS18037A).

REFERENCES

- [1] B. Biggio, B. Nelson, and P. Laskov. Support vector machines under adversarial label noise. In *Proc. of Asian Conference on Machine Learning (ACML)*, pages 97–112, 2011.
- [2] B. Biggio, I. Corona, D. Maiorca, B. Nelson, N. Šrndić, P. Laskov, G. Giacinto, and F. Roli. Evasion attacks against machine learning at test time. In *Machine Learning and Knowledge Discovery in Databases*, pages 387–402. Springer, 2013.
- [3] N. Carlini and D. A. Wagner. Towards evaluating the robustness of neural networks. In *Proc. of IEEE Symposium on Security and Privacy (S&P)*, 2017.
- [4] X. Chen, C. Liu, B. Li, K. Lu, and D. Song. Targeted backdoor attacks on deep learning systems using data poisoning. Technical report, arXiv:1712.05526, 2017.
- [5] T. Gu, B. Dolan-Gavitt, and S. Garg. Badnets: Identifying vulnerabilities in the machine learning model supply chain. Technical report, arXiv:1708.06733, 2017.
- [6] M. Kloft and P. Laskov. Online anomaly detection under adversarial impact. In *JMLR Workshop and Conference Proceedings, Volume 9: AISTATS*, pages 405–412, 2010.
- [7] A. Krizhevsky, G. Hinton, et al. Learning multiple layers of features from tiny images. Technical report, 2009.
- [8] Y. Liu, S. Ma, Y. Aafer, W.-C. Lee, J. Zhai, W. Wang, and X. Zhang. Trojaning attack on neural networks. In *Proc. of Network and Distributed System Security Symposium (NDSS)*, 2018.
- [9] E. Quiring, A. Maier, and K. Rieck. Misleading authorship attribution of source code using adversarial learning. In *Proc. of USENIX Security Symposium*, pages 479–496, 2019.
- [10] E. Quiring, D. Klein, D. Arp, M. Johns, and K. Rieck. Adversarial preprocessing: Understanding and preventing image-scaling attacks in machine learning. In *Proc. of USENIX Security Symposium*, 2020.
- [11] O. Russakovsky, J. Deng, H. Su, J. Krause, S. Satheesh, S. Ma, Z. Huang, A. Karpathy, A. Khosla, M. Bernstein, A. C. Berg, and L. Fei-Fei. ImageNet Large Scale Visual Recognition Challenge. *International Journal of Computer Vision (IJCV)*, 115(3):211–252, 2015.
- [12] A. Shafahi, W. R. Huang, M. Najibi, O. Suci, C. Studer, T. Dumitras, and T. Goldstein. Poison frogs! Targeted clean-label poisoning attacks on neural networks. In *Advances in Neural Information Processing Systems (NIPS)*, pages 6103–6113, 2018.
- [13] R. Shokri, M. Stronati, C. Song, and V. Shmatikov. Membership inference attacks against machine learning models. In *Proc. of IEEE Symposium on Security and Privacy (S&P)*, 2017.
- [14] K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. Technical report, arXiv:1409.1556, 2014.
- [15] C. Stoneff. The insider versus the outsider: Who poses the biggest security risk? <https://www.helpnetsecurity.com/2015/08/19/the-insider-versus-the-outsider-who-poses-the-biggest-security-risk/>, 2015. Accessed: 2020-01-07.
- [16] C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna. Rethinking the inception architecture for computer vision. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2818–2826, 2015.
- [17] T. J. L. Tan and R. Shokri. Bypassing backdoor detection algorithms in deep learning. Technical report, arXiv:1905.13409, 2019.
- [18] F. Tramèr, F. Zhang, A. Juels, M. K. Reiter, and T. Ristenpart. Stealing machine learning models via prediction apis. In *Proc. of USENIX Security Symposium*, pages 601–618, 2016.
- [19] B. Wang, Y. Yao, S. Shan, H. Li, B. Viswanath, H. Zheng, and B. Y. Zhao. Neural cleanse: Identifying and mitigating backdoor attacks in neural networks. In *Proc. of IEEE Symposium on Security and Privacy (S&P)*, pages 707–723, 2019.
- [20] Q. Xiao, Y. Chen, C. Shen, Y. Chen, and K. Li. Seeing is not believing: Camouflage attacks on image scaling algorithms. In *Proc. of USENIX Security Symposium*, pages 443–460, 2019.
- [21] Y. Yao, H. Li, H. Zheng, and B. Y. Zhao. Latent backdoor attacks on deep neural networks. In *Proc. of ACM Conference on Computer and Communications Security (CCS)*, pages 2041–2055, 2019.